

OPTIMIZING POSTGRESQL PERFORMANCE IN DJANGO

Ahmad Fanaei

A BIT ABOUT ME

- Python & Django developer
- backend & data team @ torob.com



WHAT CAN DEGRADE DJANGO PERFORMANCE?

- n+1 select
- Select unnecessary stuff
- Blocking migrations
- Pagination
- Slow queries

WHAT TOOLS CAN BE USED TO FIND AND FIX THESE ISSUES

- django_debug_toolbar
- assertNumQueries
- squawk
- pg_stat_statement

DJANGO ORM IS AMAZING BUT ...

ALL NON-TRIVIAL ABSTRACTIONS, TO SOME DEGREE, ARE LEAKY. ¹

- unexpected query
- not optimal queries
- blocking migrations

TOOLS

DJANGO_DEBUG_TOOLBAR 2

- [this is an amazing test](#)

DJDT



django

ASSERT NUM QUERIES

```
def test_past_question(self):
    """
    The detail view of a question with a pub_date in the past
    displays the question's text.
    """
    past_question = create_question(question_text='Past Question.', days=-5)
    url = reverse('polls:detail', args=(past_question.id,))

    with self.assertNumQueries(1):
        response = self.client.get(url)
    self.assertContains(response, past_question.question_text)
```

ASSERT NUM QUERIES

```
root@6d9a68a2769f:/app# python manage.py test
polls.tests.QuestionDetailViewTests.test_past_question
Found 1 test(s).
Creating test database for alias 'default'...
System check identified no issues (0 silenced).
F
=====
FAIL: test_past_question (polls.tests.QuestionDetailViewTests)
The detail view of a question with a pub_date in the past
-----
Traceback (most recent call last):
  File "/app/polls/tests.py", line 128, in test_past_question
    response = self.client.get(url)
  File "/usr/local/lib/python3.9/site-packages/django/test/testcases.py", line 86, in
__exit__
    self.test_case.assertEqual(
AssertionError: 2 != 1 : 2 queries executed, 1 expected
Captured queries were:
1. SELECT "polls_question"."id", "polls_question"."question_text",
"polls_question"."pub_date" FROM "polls_question" WHERE ("polls_question"."pub_date" <=
'2022-02-13T22:14:30.093869+00:00'::timestamptz AND "polls_question"."id" = 1) LIMIT 21
2. SELECT "polls_choice"."id", "polls_choice"."question_id",
"polls_choice"."choice_text", "polls_choice"."votes" FROM "polls_choice" WHERE
"polls_choice"."question_id" = 1
-----
Ran 1 test in 0.034s

FAILED (failures=1)
Destroying test database for alias 'default'...
```

SQUAWK 3

```
> squawk example.sql
example.sql:2:1: warning: prefer-text-field

2 | --
3 | -- Create model Bar
4 | --
5 | CREATE TABLE "core_bar" (
6 |     "id" serial NOT NULL PRIMARY KEY,
7 |     "alpha" varchar(100) NOT NULL
8 | );

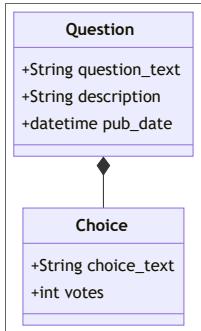
note: Changing the size of a varchar field requires an ACCESS EXCLUSIVE lock.
help: Use a text field with a check constraint.
```

PG_STAT_STATEMENT

```
postgres=# CREATE EXTENSION pg_stat_statements;
ERROR: extension "pg_stat_statements" already exists
postgres=# \d pg_stat_statements
          View "public.pg_stat_statements"
   Column    |      Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----+
  query     |    text     |           |       |       |
  calls      |    bigint    |           |       |       |
 total_time | double precision |           |       |       |
 min_time   | double precision |           |       |       |
 max_time   | double precision |           |       |       |
 mean_time  | double precision |           |       |       |
 stddev_time | double precision |           |       |       |
```

EXAMPLE

DATA MODEL



API

```
urlpatterns = [
    path('', views.IndexView.as_view(), name='index'),
    path('<int:pk>/', views.DetailView.as_view(), name='detail'),
    path('<int:pk>/results/', views.ResultsView.as_view(), name='results'),
    path('<int:question_id>/vote/', views.vote, name='vote'),
]
```

N+1 QUERY PROBLEM

Question number 1

- a1 -- 0 votes
- b1 -- 1 vote
- c1 -- 2 votes

Question number 2

- a2 -- 0 votes
- b2 -- 1 vote
- c2 -- 2 votes

Question number 3

- a3 -- 0 votes
- b3 -- 1 vote
- c3 -- 2 votes

Question number 4

- a4 -- 0 votes
- b4 -- 1 vote
- c4 -- 2 votes

Question number 5

- a5 -- 0 votes
- b5 -- 1 vote
- c5 -- 2 votes

django

```
class IndexView(generic.ListView):
    template_name = 'polls/index.html'
    context_object_name = 'latest_question_list'

    def get_queryset(self):
        """
        Return the last five published questions
        """
        return (
            Question.objects
            .filter(pub_date__lte=timezone.now())
            .order_by('-pub_date')[:5]
        )
```

```
<div style="padding: 16px;">
{%
    for question in latest_question_list %}
        <a href="{% url 'polls:detail' question.id %}">{{ question.question_text }}</a>
        <ul style="margin-top: 0">
            {% for choice in question.choice_set.all %}
                <li>{{ choice.choice_text }} -- {{ choice.votes }} vote{{ choice.votes|pluralize }}</li>
            {% endfor %}
        </ul>
    {% endfor %}
</div>
```

Question number 1

- c1 -- 2 votes
- b1 -- 1 vote
- a1 -- 0 votes

Question number 2

- c2 -- 2 votes
- b2 -- 1 vote
- a2 -- 0 votes

Question number 3

- c3 -- 2 votes
- b3 -- 1 vote
- a3 -- 0 votes

Question number 4

- c4 -- 2 votes
- b4 -- 1 vote
- a4 -- 0 votes

Question number 5

- c5 -- 2 votes
- b5 -- 1 vote
- a5 -- 0 votes

django

Powered by

`select_related`: Returns a QuerySet that will "follow" foreign-key relationships

`prefetch_related`: Returns a QuerySet that will automatically retrieve, in a single batch, related objects

```
class IndexView(generic.ListView):
    template_name = 'polls/index.html'
    context_object_name = 'latest_question_list'

    def get_queryset(self):
        """
        Return the last five published questions
        """
        return (
            Question.objects
            .prefetch_related('choice_set')
            .filter(pub_date__lte=timezone.now())
            .order_by('-pub_date')[:5]
        )
```

Question number 1

- a1 -- 0 votes
- b1 -- 1 vote
- c1 -- 2 votes

Question number 2

- a2 -- 0 votes
- b2 -- 1 vote
- c2 -- 2 votes

Question number 3

- a3 -- 0 votes
- b3 -- 1 vote
- c3 -- 2 votes

Question number 4

- a4 -- 0 votes
- b4 -- 1 vote
- c4 -- 2 votes

Question number 5

- a5 -- 0 votes
- b5 -- 1 vote
- c5 -- 2 votes

django

ASSERT NUM QUERIES

```
def test_past_question(self):
    """
    Questions with a pub_date in the past are displayed on the
    index page.
    """
    create_question(question_text="Past question.", days=-30)
    create_question(question_text="Past question. 2", days=-15)
    with self.assertNumQueries(2):
        self.client.get(reverse('polls:index'))
```

EXCESS SELECTION

By default django selects all columns.

```
list(  
    Question.objects  
        .prefetch_related('choice_set')  
        .filter(pub_date_lte=timezone.now())  
        .order_by('-pub_date')[:5]  
)
```

```
SELECT
    "polls_question"."id",
    "polls_question"."question_text",
    "polls_question"."description",
    "polls_question"."pub_date"
FROM "polls_question"
WHERE
    "polls_question"."pub_date" <= \'2022-02-14T21:19:25.545889+00:00\'::timestamptz
ORDER BY "polls_question"."pub_date"
DESC LIMIT 5
```

Limit selected columns using values or only

```
list(
    Question.objects
        .prefetch_related('choice_set')
        .only('question_text')
        .filter(pub_date__lte=timezone.now())
        .order_by('-pub_date')[:5]
)
```

```
SELECT
    "polls_question"."id",
    "polls_question"."question_text"
FROM "polls_question"
WHERE "polls_question"."pub_date" <= \'2022-02-14T21:21:36.022491+00:00
\':timestamptz
ORDER BY "polls_question"."pub_date"
DESC LIMIT 5
```

A FEW MONTH LATER

[Question 2](#) 2022-02-13

- c4 -- 2 votes
- b4 -- 1 vote
- a4 -- 0 votes



[Question 3](#) 2022-02-13

- c7 -- 2 votes
- b7 -- 1 vote
- a7 -- 0 votes

[Question 4](#) 2022-02-13

- c10 -- 2 votes
- b10 -- 1 vote
- a10 -- 0 votes

[Question 5](#) 2022-02-13

- c13 -- 2 votes
- b13 -- 1 vote
- a13 -- 0 votes

[Question 1](#) 2022-02-13

- c1 -- 2 votes
- b1 -- 1 vote
- a1 -- 0 votes

django

ALWAYS PREFER values OVER only

BLOCKING MIGRATIONS

MAKE QUESTIONS UNIQUE

```
question_text = models.CharField(max_length=200, unique=True)
```

```
BEGIN;
--
-- Alter field question_text on question
--
ALTER TABLE "polls_question" ADD CONSTRAINT
"polls_question_question_text_e0c682c7_uniq" UNIQUE ("question_text");
CREATE INDEX "polls_question_question_text_e0c682c7_like" ON "polls_question"
("question_text" varchar_pattern_ops);
COMMIT;
```

... a standard index build locks out writes (but not reads) on the table until it's done

SQUAWK

```
root@6d9a68a2769f:/app# ./squawk-linux-x86_64 out.sql
out.sql:2:1: warning: disallowed-unique-constraint

2 | --
3 | -- Alter field question_text on question
4 | --
5 | ALTER TABLE "polls_question" ADD CONSTRAINT
"polls_question_question_text_e0c682c7_uniq" UNIQUE ("question_text");

note: Adding a UNIQUE constraint requires an ACCESS EXCLUSIVE lock which blocks
reads.
  help: Create an index CONCURRENTLY and create the constraint using the index.
```

```
out.sql:6:1: warning: require-concurrent-index-creation

  6 | CREATE INDEX "polls_question_question_text_e0c682c7_like" ON "polls_question"
("question_text" varchar_pattern_ops);

  note: Creating an index blocks writes.
  help: Create the index CONCURRENTLY.

find detailed examples and solutions for each rule at https://squawkhq.com/docs/rules
```

PAGINATION

WHAT IS THE PROBLEM

[Question 10005](#) 2022-02-13

- c30013 -- 2 votes
- b30013 -- 1 vote
- a30013 -- 0 votes



[Question 10004](#) 2022-02-13

- c30010 -- 2 votes
- b30010 -- 1 vote
- a30010 -- 0 votes

[Question 10003](#) 2022-02-13

- c30007 -- 2 votes
- b30007 -- 1 vote
- a30007 -- 0 votes

[Question 10002](#) 2022-02-13

- c30004 -- 2 votes
- b30004 -- 1 vote
- a30004 -- 0 votes

[Question 10001](#) 2022-02-13

- c30001 -- 2 votes
- b30001 -- 1 vote
- a30001 -- 0 votes

[« Previous](#) [Next »](#)

django

DEFAULT PAGINATION

```
class IndexView(generic.ListView):
    template_name = 'polls/index.html'
    context_object_name = 'latest_question_list'
    paginate_by = 5

    def get_queryset(self):
        """
        Return the last five published questions (not including those set to be
        published in the future).
        """
        return Question.objects.prefetch_related('choice_set').filter(
            pub_date__lte=timezone.now()
        ).order_by('-pub_date', '-id')
```

```
<div class="pagination">
    <span class="previous">
        {%
            if page_obj.has_previous %}
            <a href="?page={{ page_obj.previous_page_number }} ">&laquo;
        Previous</a>
        {%
            endif %}
    </span>
    <span class="next">
        {%
            if page_obj.has_next %}
            <a href="?page={{ page_obj.next_page_number }} ">Next &raquo;</a>
        {%
            endif %}
    </span>
</div>
```

[Question 20000](#) 2022-02-13

- c59998 -- 2 votes
- b59998 -- 1 vote
- a59998 -- 0 votes



[Question 19999](#) 2022-02-13

- c59995 -- 2 votes
- b59995 -- 1 vote
- a59995 -- 0 votes

[Question 19998](#) 2022-02-13

- c59992 -- 2 votes
- b59992 -- 1 vote
- a59992 -- 0 votes

[Question 19997](#) 2022-02-13

- c59989 -- 2 votes
- b59989 -- 1 vote
- a59989 -- 0 votes

[Question 19996](#) 2022-02-13

- c59986 -- 2 votes
- b59986 -- 1 vote
- a59986 -- 0 votes

[Next >](#)

django

WHY IS DEFAULT PAGINATION SLOW

- Counting number of items. **why?**
 - wasn't it possible to store count separately?
 - we have index counting number of items using an index should be trivial!
- offset-limit pagination becomes slower in high pages

LIMIT-OFFSET

```
SELECT *
FROM "polls_question"
WHERE "polls_question"."pub_date" <= '2022-02-14T22:07:52.796083+00:00'::timestamptz
ORDER BY "polls_question"."id" DESC
LIMIT 5
OFFSET 9995
```

KEYSET

```
SELECT *
  FROM "polls_question"
 WHERE
    "polls_question"."pub_date" <= '2022-02-14T22:07:52.796083+00:00'::timestamptz
      AND "polls_question"."id" <= ...
ORDER BY "polls_question"."id" DESC
LIMIT 5
```

DJANGO-KEYSET-PAGINATION-PLUS

KEYSET PAGINATION ISSUES

- User can't see total number of pages
- User can't paginate to a specific page and must skip pages one-by-one

SLOW QUERY

FINDING SLOW QUERIES

```
SELECT
    sum(total_time) / 1000 / 60 as total_time,
    sum(calls) as total_calls,
    query
FROM pg_stat_statements
GROUP BY query
ORDER BY 1 DESC
LIMIT 200;
```

SAMPLE OUTPUT

total_time	total_calls	query
0.01313467055	12	SELECT "polls_question"."id", "polls_question"."question_text", "polls_question"."description", "polls_question"."pub_date" FROM "polls_question" WHERE "polls_question"."pub_date" <= \$1::timestamptz ORDER BY "polls_question"."pub_date" DESC, "polls_question"."id" DESC LIMIT \$2 OFFSET \$3
0.000491445183333333	12	SELECT COUNT(*) AS "__count" FROM "polls_question" WHERE "polls_question"."pub_date" <= \$1::timestamptz
1.65117e-05	12	SELECT "auth_user"."id", "auth_user"."password", "auth_user"."last_login", "auth_user"."is_superuser", "auth_user"."username", "auth_user"."first_name", "auth_user"."last_name", "auth_user"."email", "auth_user"."is_staff", "auth_user"."is_active", "auth_user"."date_joined" FROM "auth_user" WHERE "auth_user"."id" = \$1 LIMIT \$2
1.452433333333333e-05	12	SELECT "polls_choice"."id", "polls_choice"."question_id", "polls_choice"."choice_text", "polls_choice"."votes" FROM "polls_choice" WHERE "polls_choice"."question_id" IN (\$1, \$2, \$3, \$4, \$5)
1.184633333333333e-05	12	SELECT "django_session"."session_key", "django_session"."session_data", "django_session"."expire_date" FROM "django_session" WHERE ("django_session"."expire_date" > \$1::timestamptz AND "django_session"."session_key" = \$2) LIMIT \$3

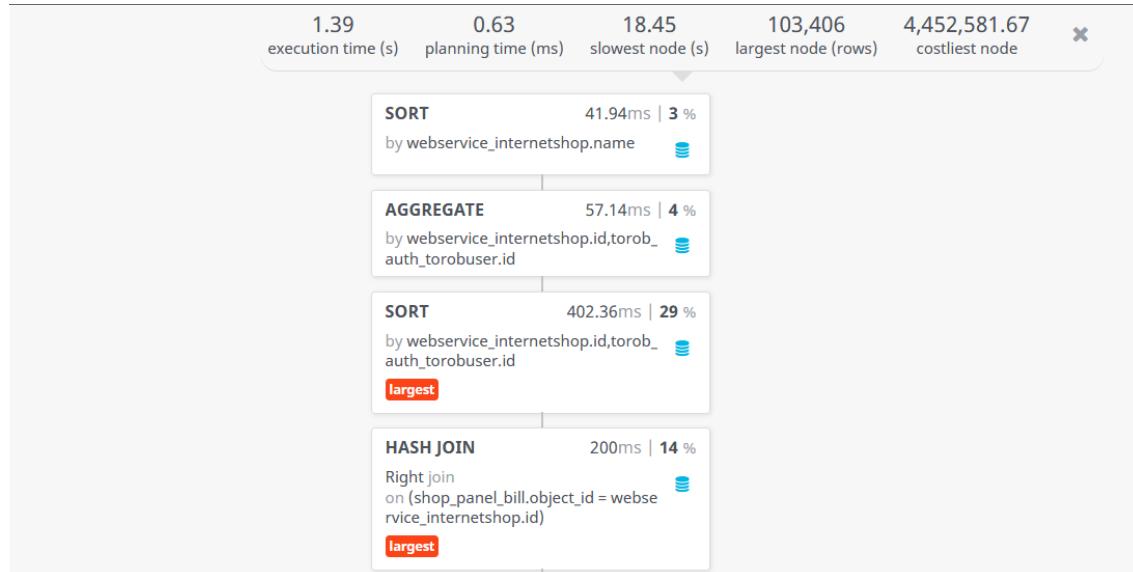
EXPLAIN OUTPUT

```
SELECT "polls_choice"."id",
       "polls_choice"."question_id",
       "polls_choice"."choice_text",
       "polls_choice"."votes"
  FROM "polls_choice"
 WHERE "polls_choice"."question_id" IN (202922, 202921, 202920, 202919, 202918)
```

EXPLAIN OUTPUT

```
QUERY PLAN
Index Scan using polls_choice_question_id_c5b4b260 on polls_choice  (cost=0.29..25.71
rows=15 width=18) (actual time=0.056..0.063 rows=15 loops=1)
  Index Cond: (question_id = ANY ('{202922,202921,202920,202919,202918}'::integer[]))
Planning Time: 1.078 ms
Execution Time: 0.139 ms
```

VISUALIZE EXPLAIN OUT PUT 4



HOW TO MAKE QUERIES FASTER

- add index or partial index
- increase statics collection
- tune auto-vaccum

THANK YOU

Premature optimization is the root of all evil – Donald Knuth